


# NumPy and Pandas Fundamentals: A Step-by-Step Guide for Data Analysis

Wei-Hsin Hsu, 2023/12/9

 [Linkedin](#)

 [Github](#)



In this document, we will cover fundamental techniques of NumPy and Pandas, two open-source Python libraries that are very helpful for data manipulation, data analysis, and multi-dimensional data processing. They provide convenient and practical data structures and functions, and are widely used in the fields of data analysis.

Once you master the techniques mentioned in this document, you will have the ability to perform basic data processing, analysis, and visualization, taking the first step towards becoming a data analyst. :)

## Table of Contents

<b>NumPy .....</b>	<b>1</b>
<b>1. NumPy Array</b>	
- Creating, Datatype, Basic Operations	
- 2D Array, Shape, Transpose, Creating Array with 0 and 1	
<b>2. Sequence</b>	
- Arithmetic Sequence with Specified intervals or Number of Elements	
<b>3. Operation Functions</b>	
- Statistical, Mathematical Operations	
<b>4. Creating Array with Random Values</b>	
<b>5. Selecting Data in NumPy Array</b>	
- Changing Values, Filtering and more Operations	
<b>6. Reshaping NumPy Arrays</b>	
<b>7. Concatenate NumPy Arrays</b>	
<b>8. Changing Datatype</b>	
<b>9. Sorting</b>	
<b>10. Splitting Data</b>	
- Horizontal, Vertical Splitting	
<b>Pandas .....</b>	<b>9</b>
<b>1. Series</b>	
- Creating, Datatype, Index	
- Selecting Data	
<b>2. DataFrame</b>	
- Creating, Datatype, Importing Data and Saving Data	
- Selecting and Dropping a Column	
- Mathematical Operations, Check Data Distribution	
- Filtering, Sorting Data	
<b>3. Working with Missing Data</b>	
- Checking, Replacing Missing Values	
<b>4. Working with Duplicate Data</b>	
- Checking, Dropping Duplicate Values	
<b>5. Renaming a Column</b>	
<b>6. Grouping Data</b>	
- Groupby, Aggregate	
<b>7. Data Overview</b>	
- Info, Changing Datatype, Describe	
<b>8. Basic Visualization with Pandas</b>	
<b>Conclusion .....</b>	<b>24</b>

# NumPy

NumPy (short for Numerical Python) is a powerful numerical computing library in Python. It offers many mathematical functions and operations, particularly tailored for multi-dimensional arrays. It provides scientists, engineers, and data scientists with rich tools, making calculations and data analysis in Python more easily and efficient.

We load the NumPy library and commonly use 'np' as an alias.

```
In [1]: import numpy as np
```

## 1. NumPy Array

### Creating a NumPy Array

Elements of a NumPy array can be assigned using a list or tuple.

```
In [2]: A= np.array([1, 2, 3])
        B= np.array((4, 5, 6))
        print(A)
        print(B)
```

```
[1 2 3]
[4 5 6]
```

### Datatype of the NumPy Array

```
In [3]: type(A)
```

```
Out[3]: numpy.ndarray
```

### Basic Operations

```
In [4]: print(A**2)
```

```
[1 4 9]
```

```
In [5]: print(A+B)
        print(A-B)
        print(A*B)
        print(A/B)
```

```
[5 7 9]
[-3 -3 -3]
[ 4 10 18]
[0.25 0.4 0.5 ]
```

## 2-dimensional NumPy Array

```
In [6]: C = np.array([[1, 2, 3], [4, 5, 6]])  
D = np.array([[7, 8, 9], [10, 11, 12]])  
print(C)  
print(D)
```

```
[[1 2 3]  
 [4 5 6]]  
[[ 7  8  9]  
 [10 11 12]]
```

For NumPy arrays with more dimensions, you just need to add data of the same format accordingly.

```
In [7]: print(C+D)  
# print(C-D)  
# print(C*D)  
# print(C/D)
```

```
[[ 8 10 12]  
 [14 16 18]]
```

## Shape, Size, and Dimension

```
In [8]: A = np.array([[1, 2, 3, 4, 5], [6, 7, 8, 9, 10]])  
B = np.array([3, 2, 1])  
  
#Size  
print(A.shape)  
print(B.shape)
```

```
(2, 5)  
(3,)
```

```
In [9]: #Values count  
print(A.size)  
print(B.size)
```

```
10  
3
```

```
In [10]: #Dimension  
print(A.ndim)  
print(B.ndim)
```

```
2  
1
```

## Transpose

Transposing a 2-D array swaps its rows and columns.

```
In [11]: A = np.array([[1,2,3],[4,5,6]])  
A.transpose() # This is equivalent to A.T
```

```
Out[11]: array([[1, 4],  
               [2, 5],  
               [3, 6]])
```

ravel() can convert multi-dimensional array into 1D array.

```
In [12]: A.ravel()
```

```
Out[12]: array([1, 2, 3, 4, 5, 6])
```

## Creating Array with 0 and 1

np.zeros() creates an array of specified size with 0, while np.ones() creates an array of specified size with 1.

```
In [13]: print(np.zeros(10))  
print(np.ones(10))
```

```
[0. 0. 0. 0. 0. 0. 0. 0. 0. 0.]  
[1. 1. 1. 1. 1. 1. 1. 1. 1. 1.]
```

```
In [14]: np.zeros((2,4))
```

```
Out[14]: array([[0., 0., 0., 0.],  
               [0., 0., 0., 0.]])
```

```
In [15]: np.ones((3,2,1))
```

```
Out[15]: array([[[1.],  
                 [1.]],  
               [[1.],  
                 [1.]],  
               [[1.],  
                 [1.]])
```

## 2. Sequence

### Arithmetic Sequence with Specified Intervals

```
In [16]: np.arange(0,100,25)
```

```
Out[16]: array([ 0, 25, 50, 75])
```

### Arithmetic Sequence with a Specified Number of Elements.

Noted that argument 'endpoint = True' meaning includes stop, and vice versa.(Default is True)

```
In [17]: np.linspace(0,100,26, endpoint= True)
```

```
Out[17]: array([ 0.,  4.,  8., 12., 16., 20., 24., 28., 32., 36., 40.,
                44., 48., 52., 56., 60., 64., 68., 72., 76., 80., 84.,
                88., 92., 96., 100.])
```

## 3. Operation Functions

### Statistical Operation

```
In [18]: A = np.array([1, 3, 5, 7, 9])
print(np.mean(A))
print(np.median(A))
print(np.std(A))
print(np.var(A))
```

```
5.0
5.0
2.8284271247461903
8.0
```

### Mathematical Operation

```
In [19]: num= 100
print(np.sqrt(num)) #Square root
print(np.exp(num)) #Exponential
print(np.log(num)) #Logarithm with the base e
print(np.log10(num)) #Logarithm with the base 10
#print(np.sin(num))
#print(np.cos(num))
#.....
```

```
10.0
2.6881171418161356e+43
4.605170185988092
2.0
```

Calculating Inner product (Dot).

```
In [20]: A= np.array([1, 2, 3])
B= np.array([4, 5, 6])
C= np.array([1, 4, 7])
print(np.dot(A, B))
print(np.dot(A, C))
```

```
32
30
```

## 4. Creating Array with Random Values

`random.randn()` is used to randomly generate data based on the Standard Normal Distribution.

```
In [21]: np.random.randn(5)
```

```
Out[21]: array([ 1.6485286 , -1.46597283, -0.37945557, -0.26006813, -0.10395762])
```

```
In [22]: np.random.randn(2, 3)
```

```
Out[22]: array([[ -0.03435319,  0.71343253, -0.71856028],  
                [-0.14204424, -0.2228869 , -3.11251075]])
```

`numpy.random.randint(low=x, high=y, size=z)` generates `z` integers from `x` to `y`, including `x` but excluding `y`.

```
In [23]: np.random.randint(low= 1, high= 10, size=10)
```

```
Out[23]: array([7, 7, 1, 8, 5, 6, 2, 7, 5, 5])
```

## 5. Selecting Data in NumPy Array

```
In [24]: A= np.array([1, 2, 3, 4, 5, 6])  
        B= np.array((7, 8, 9, 10, 11, 12))  
        A[0]
```

```
Out[24]: 1
```

```
In [25]: A[2:4]
```

```
Out[25]: array([3, 4])
```

### Changing Values

```
In [26]: A[2]= 11  
        B[0]= 3  
        print(A)  
        print(B)
```

```
[ 1  2 11  4  5  6]  
[ 3  8  9 10 11 12]
```

## Filtering and more Operation

```
In [27]: C = np.array([[1, 2, 3], [4, 5, 6]])
D = np.array([[7, 8, 9], [10, 11, 12]])

print(C[0,1])
print(D[0,0:2])
```

```
2
[7 8]
```

```
In [28]: print(C>2)
print(C[C>2])
```

```
[[False False  True]
 [ True  True  True]]
[3 4 5 6]
```

```
In [29]: print(C.min())
print(np.max(C, axis=1)) # This is equivalent to C.max(axis=1)
#np.sum()
#np.product()
#.....
```

```
1
[3 6]
```

```
In [30]: print(C.sum())
print(C.sum(axis=0))
print(C.mean(axis=1))
```

```
21
[5 7 9]
[2. 5.]
```

## 6. Reshaping NumPy Arrays

```
In [31]: #Reshape
A = np.array([1,2,3,4,5,6,7,8])
A.reshape(2,4)
```

```
Out[31]: array([[1, 2, 3, 4],
               [5, 6, 7, 8]])
```

## 7. Concatenate NumPy Arrays

```
In [32]: A = np.array([1, 2, 3, 4, 5, 6, 7, 8])
B = np.array([7, 6, 5, 4, 3, 2, 1, 0])
print(np.concatenate((A,B)))
```

```
[1 2 3 4 5 6 7 8 7 6 5 4 3 2 1 0]
```



## 8. Changing Datatype

```
In [33]: A= np.array([1, 2, 3])
print(A.dtype)
B= np.array([[1.5, 2.5, 3.5],[4, 5, 6]])
print(B.dtype)
```

```
int32
float64
```

```
In [34]: C= A.astype(np.float64)
print(C.dtype)
D= B.astype(np.int64)
print(D)
print(D.dtype)
```

```
float64
[[1 2 3]
 [4 5 6]]
int64
```

## 9. Sorting Data

```
In [35]: A= np.array([3, 2, 5, 4, 1])
np.sort(A)
```

```
Out[35]: array([1, 2, 3, 4, 5])
```

```
In [36]: A= np.array([[1,3,2],[4,6,5]])
print(np.sort(A))
print(np.sort(A, axis= 0))
```

```
[[1 2 3]
 [4 5 6]]
[[1 3 2]
 [4 6 5]]
```

## 10. Splitting Data

```
In [37]: A= np.arange(16).reshape(4,4)
A
```

```
Out[37]: array([[ 0,  1,  2,  3],
                [ 4,  5,  6,  7],
                [ 8,  9, 10, 11],
                [12, 13, 14, 15]])
```

## Horizontal Splitting

```
In [38]: np.hsplit(A,2) #Splitting to 2 equal parts
```

```
Out[38]: [array([[ 0,  1],
                [ 4,  5],
                [ 8,  9],
                [12, 13]]),
          array([[ 2,  3],
                [ 6,  7],
                [10, 11],
                [14, 15]])]
```

```
In [39]: np.hsplit(A,4) #Splitting to 4 equal parts
```

```
Out[39]: [array([[ 0],
                [ 4],
                [ 8],
                [12]]),
          array([[ 1],
                [ 5],
                [ 9],
                [13]]),
          array([[ 2],
                [ 6],
                [10],
                [14]]),
          array([[ 3],
                [ 7],
                [11],
                [15]])]
```

## Vertical Splitting

```
In [40]: np.vsplit(A,4) #Splitting to 4 equal parts
```

```
Out[40]: [array([[0, 1, 2, 3]]),
          array([[4, 5, 6, 7]]),
          array([[ 8,  9, 10, 11]]),
          array([[12, 13, 14, 15]])]
```

# Pandas

Pandas (derived from "Panel Data") is a powerful library in Python that provides high-performance, user-friendly data structures and functions for data analysis. Pandas serves as an ideal tool for handling structured data, making data cleaning, processing, and analysis more accessible in Python.

We load the Pandas library and commonly use 'pd' as an alias.

```
In [41]: import pandas as pd
```

## 1. Series

A Series is a one-dimensional array of data, much like a column in Excel, and each data has an index.

### Creating a Series

```
In [42]: data_series = pd.Series([1, 2, 3, 4, 5])
data_series
```

```
Out[42]: 0    1
         1    2
         2    3
         3    4
         4    5
         dtype: int64
```

### Datatype of the Series

```
In [43]: type(data_series)
```

```
Out[43]: pandas.core.series.Series
```

### Index

```
In [44]: data_series.index
```

```
Out[44]: RangeIndex(start=0, stop=5, step=1)
```

We did not assign an index for the series at first, so the index defaults to numbers starting from 0.

### Values

```
In [45]: data_series.values
```

```
Out[45]: array([1, 2, 3, 4, 5], dtype=int64)
```

## Selecting Data in the Series

Let's use stock price data for 3 companies at a certain point of time as an example.

```
In [46]: st= {"AAPL":194.27, "META": 326.59, "NVDA":465.96}
st_series= pd.Series(st)
st_series
```

```
Out[46]: AAPL    194.27
         META    326.59
         NVDA    465.96
         dtype: float64
```

Select the data in the 2nd place, noticed that Python uses zero-based indexing.

```
In [47]: st_series[1]
```

```
Out[47]: 326.59
```

Select multiple data.

```
In [48]: st_series[[0,2]]
```

```
Out[48]: AAPL    194.27
         NVDA    465.96
         dtype: float64
```

Use the Index to select data.

```
In [49]: st_series["AAPL"]
```

```
Out[49]: 194.27
```

## 2. Dataframe

A DataFrame is a two-dimensional data, it has rows and columns in a table format and also has an index for each observation.

### Creating a DataFrame

Let's use the daily stock price data for the same companies over a period of 10 days.

※Noted that this data is fictional, and in practice, stock price data may contain missing values because the stock market is closed during holidays.

```
In [50]: st_m = {"Date":["2023-12-08", "2023-12-09", "2023-12-10", "2023-12-11", "2023-12-12", "2023-12-13", "2023-12-14", "2023-12-15", "2023-12-16", "2023-12-17"],
               "AAPL": [194.27, 198.32, 199.56, 201.19, 200.35, 201.88, 202.52, 201.47, 203.51, 204.76],
               "META": [326.59, 330.12, 327.31, 339.51, 336.87, 338.12, 340.53, 339.12, 342.90, 345.19],
               "NVDA": [465.97, 467.13, 471.52, 466.31, 465.19, 466.12, 470.67, 468.53, 469.09, 467.13]}
df = pd.DataFrame(st_m)
df
```

Out[50]:

	Date	AAPL	META	NVDA
0	2023-12-08	194.27	326.59	465.97
1	2023-12-09	198.32	330.12	467.13
2	2023-12-10	199.56	327.31	471.52
3	2023-12-11	201.19	339.51	466.31
4	2023-12-12	200.35	336.87	465.19
5	2023-12-13	201.88	338.12	466.12
6	2023-12-14	202.52	340.53	470.67
7	2023-12-15	201.47	339.12	468.53
8	2023-12-16	203.51	342.90	469.09
9	2023-12-17	204.76	345.19	467.13

### Datatype of the DataFrame

```
In [51]: type(df)
```

Out[51]: pandas.core.frame.DataFrame

## Other Method of Importing Data

If the data is stored in another format, it can be loaded using `pd.read_filetype("file")`, such as Excel, CSV, etc.

```
In [52]: # df= pd.read_excel("file")
# df= pd.read_csv("file")
# .....
```

## Saving the Data as Other Format

Once you complete programming, you can save the file as excel, csv, etc. by using these code.

```
In [53]: # df.to_excel("file")
# df.to_csv("file")
# .....
```

## Selecting Data in the DataFrame

```
In [54]: df.loc[1]
```

```
Out[54]: Date      2023-12-09
AAPL          198.32
META          330.12
NVDA          467.13
Name: 1, dtype: object
```

```
In [55]: df.iloc[1]
```

```
Out[55]: Date      2023-12-09
AAPL          198.32
META          330.12
NVDA          467.13
Name: 1, dtype: object
```

It might appear that `loc` and `iloc` are the same, but they are not.

`loc` is used to retrieve values based on labels, while `iloc` is used to retrieve values based on integer positions of columns.

## Dropping a Column

```
In [56]: df= pd.DataFrame(st_m, index= st_m["Date"] ) #Set index = Date
df.drop(["Date"],axis=1, inplace= True) #Because We've set index to Date, s
df
```

Out[56]:

	AAPL	META	NVDA
2023-12-08	194.27	326.59	465.97
2023-12-09	198.32	330.12	467.13
2023-12-10	199.56	327.31	471.52
2023-12-11	201.19	339.51	466.31
2023-12-12	200.35	336.87	465.19
2023-12-13	201.88	338.12	466.12
2023-12-14	202.52	340.53	470.67
2023-12-15	201.47	339.12	468.53
2023-12-16	203.51	342.90	469.09
2023-12-17	204.76	345.19	467.13

Noticed that the 'inplace' argument above.

When working with data, if you want to make modifications to the original object, you need to set 'inplace =True'. Otherwise, the changes will not be applied to the original object.

You can see that loc can be used to select data based on index labels.

```
In [57]: df.loc["2023-12-11"]
```

Out[57]: AAPL 201.19  
META 339.51  
NVDA 466.31  
Name: 2023-12-11, dtype: float64

## Selecting a Column

```
In [58]: df["AAPL"]
```

Out[58]: 2023-12-08 194.27  
2023-12-09 198.32  
2023-12-10 199.56  
2023-12-11 201.19  
2023-12-12 200.35  
2023-12-13 201.88  
2023-12-14 202.52  
2023-12-15 201.47  
2023-12-16 203.51  
2023-12-17 204.76  
Name: AAPL, dtype: float64

## Mathematical Operations on an Column

```
In [59]: df["AAPL"].mean()  
#df["AAPL"].sum()  
#df["AAPL"].max()  
#.....
```

```
Out[59]: 200.78300000000002
```

## Check the Data Distribution

```
In [60]: df["AAPL"].value_counts()
```

```
Out[60]: 194.27    1  
198.32    1  
199.56    1  
201.19    1  
200.35    1  
201.88    1  
202.52    1  
201.47    1  
203.51    1  
204.76    1  
Name: AAPL, dtype: int64
```

## Filtering Data

```
In [61]: df["AAPL"] > 200
```

```
Out[61]: 2023-12-08    False  
2023-12-09    False  
2023-12-10    False  
2023-12-11     True  
2023-12-12     True  
2023-12-13     True  
2023-12-14     True  
2023-12-15     True  
2023-12-16     True  
2023-12-17     True  
Name: AAPL, dtype: bool
```

Filtering Data with multiple conditions.

```
In [62]: df[(df["AAPL"] > 200) & (df["NVDA"] > 470)]
```

```
Out[62]:
```

	AAPL	META	NVDA
2023-12-14	202.52	340.53	470.67



## Sorting Data

```
In [63]: df.sort_values(by = "AAPL")
```

Out[63]:

	AAPL	META	NVDA
2023-12-08	194.27	326.59	465.97
2023-12-09	198.32	330.12	467.13
2023-12-10	199.56	327.31	471.52
2023-12-12	200.35	336.87	465.19
2023-12-11	201.19	339.51	466.31
2023-12-15	201.47	339.12	468.53
2023-12-13	201.88	338.12	466.12
2023-12-14	202.52	340.53	470.67
2023-12-16	203.51	342.90	469.09
2023-12-17	204.76	345.19	467.13

Sorting data in descending order.

```
In [64]: df.sort_values(by= "AAPL", ascending= False)
```

Out[64]:

	AAPL	META	NVDA
2023-12-17	204.76	345.19	467.13
2023-12-16	203.51	342.90	469.09
2023-12-14	202.52	340.53	470.67
2023-12-13	201.88	338.12	466.12
2023-12-15	201.47	339.12	468.53
2023-12-11	201.19	339.51	466.31
2023-12-12	200.35	336.87	465.19
2023-12-10	199.56	327.31	471.52
2023-12-09	198.32	330.12	467.13
2023-12-08	194.27	326.59	465.97

### Checking the Data in the First n Row.

For the last n row, you can use `.tail(n)` instead.

```
In [65]: df.head()  
#df.tail()
```

```
Out[65]:
```

	AAPL	META	NVDA
2023-12-08	194.27	326.59	465.97
2023-12-09	198.32	330.12	467.13
2023-12-10	199.56	327.31	471.52
2023-12-11	201.19	339.51	466.31
2023-12-12	200.35	336.87	465.19

```
In [66]: df.head(3)  
#df.tail(3)
```

```
Out[66]:
```

	AAPL	META	NVDA
2023-12-08	194.27	326.59	465.97
2023-12-09	198.32	330.12	467.13
2023-12-10	199.56	327.31	471.52

## 3. Working with Missing Data

### Checking Missing Values

```
In [67]: df.isnull()
```

```
Out[67]:
```

	AAPL	META	NVDA
2023-12-08	False	False	False
2023-12-09	False	False	False
2023-12-10	False	False	False
2023-12-11	False	False	False
2023-12-12	False	False	False
2023-12-13	False	False	False
2023-12-14	False	False	False
2023-12-15	False	False	False
2023-12-16	False	False	False
2023-12-17	False	False	False

When dealing with a large amount of data, the above method may seem a bit clumsy. We can use `sum()` to quickly check the number of missing values.

```
In [68]: df.isnull().sum()
```

```
Out[68]: AAPL      0  
         META      0  
         NVDA      0  
         dtype: int64
```

Let's use a new dataset, a DataFrame containing product names, prices, and quantities but contains some issues.

```
In [69]: from numpy import nan  
data = {"Product":["A", "B", "C", "D", "E", "F", "F", "G", "H"],  
        "Price": [25, 50, 25, 25, nan, 40, 40, 25, 15],  
        "Quantity": [4, 2, nan, 9, 4, 6, 6, 4, nan] ,  
        "GP%" : [0.2, 0.15, 0.2, 0.15, 0.4, 0.15 ,0.2, 0.5, 0.2]  
}  
data = pd.DataFrame(data, index=data["Product"])  
data.drop("Product", axis=1, inplace=True)  
data
```

```
Out[69]:
```

	Price	Quantity	GP%
A	25.0	4.0	0.20
B	50.0	2.0	0.15
C	25.0	NaN	0.20
D	25.0	9.0	0.15
E	NaN	4.0	0.40
F	40.0	6.0	0.15
F	40.0	6.0	0.20
G	25.0	4.0	0.50
H	15.0	NaN	0.20

We can see that there are some NaN values.  
dropna() will delete rows that contain missing values.

```
In [70]: data.dropna() #Noticed that we did not specified the argument inplace = True
```

```
Out[70]:
```

	Price	Quantity	GP%
A	25.0	4.0	0.20
B	50.0	2.0	0.15
D	25.0	9.0	0.15
F	40.0	6.0	0.15
F	40.0	6.0	0.20
G	25.0	4.0	0.50

If we set the argument how = "all", it will only delete rows where all values are missing.

```
In [71]: data.dropna(how="all")
```

```
Out[71]:
```

	Price	Quantity	GP%
A	25.0	4.0	0.20
B	50.0	2.0	0.15
C	25.0	NaN	0.20
D	25.0	9.0	0.15
E	NaN	4.0	0.40
F	40.0	6.0	0.15
F	40.0	6.0	0.20
G	25.0	4.0	0.50
H	15.0	NaN	0.20

### Replacing Missing Values

```
In [72]: data.fillna(0)
```

```
Out[72]:
```

	Price	Quantity	GP%
A	25.0	4.0	0.20
B	50.0	2.0	0.15
C	25.0	0.0	0.20
D	25.0	9.0	0.15
E	0.0	4.0	0.40
F	40.0	6.0	0.15
F	40.0	6.0	0.20
G	25.0	4.0	0.50
H	15.0	0.0	0.20

In practice, we often use specific values to replace missing values, such as the mean. Of course, it depends on the type of data.

In the Quantity column, we use the mean to replace missing values. As for the Price column, we use 15.

```
In [73]: data["Quantity"].fillna(data["Quantity"].mean(), inplace= True)
data["Price"].fillna(15, inplace= True)
data
```

Out[73]:

	Price	Quantity	GP%
A	25.0	4.0	0.20
B	50.0	2.0	0.15
C	25.0	5.0	0.20
D	25.0	9.0	0.15
E	15.0	4.0	0.40
F	40.0	6.0	0.15
F	40.0	6.0	0.20
G	25.0	4.0	0.50
H	15.0	5.0	0.20

## 4. Working with Duplicate Data

### Checking Duplicate Values

```
In [74]: duplicate_values = data.duplicated()
print(duplicate_values)
```

```
A    False
B    False
C    False
D    False
E    False
F    False
F    False
G    False
H    False
dtype: bool
```

The default method of the `duplicated()` is to return True if the entire row is duplicated, and it will not return True if only part of it is duplicated.

We can use the `index.duplicated()` to check if there are duplicate values in the product columns(Index).

```
In [75]: duplicate_index = data.index.duplicated()
print(f"Index has duplicates: {duplicate_index}")
```

```
Index has duplicates: [False False False False False False  True False False]
```

You can add some code to meet your specific needs, such as `any()`, `sum()`, etc.

```
In [76]: for i in data.columns:
        duplicate_values = data[i].duplicated().sum()
        print(f"Column '{i}' has duplicates: {duplicate_values}")
```

```
Column 'Price' has duplicates: 5
Column 'Quantity' has duplicates: 4
Column 'GP%' has duplicates: 5
```

## Dropping Duplicate Values

```
In [77]: data.drop_duplicates(inplace=True)
data
#data.drop_duplicates(keep="last") ## Keep only the last duplicate values.
#df.drop_duplicates(keep="False") ## Delete all duplicate values.
```

Out[77]:

	Price	Quantity	GP%
A	25.0	4.0	0.20
B	50.0	2.0	0.15
C	25.0	5.0	0.20
D	25.0	9.0	0.15
E	15.0	4.0	0.40
F	40.0	6.0	0.15
F	40.0	6.0	0.20
G	25.0	4.0	0.50
H	15.0	5.0	0.20

## 5. Renaming a Column

```
In [78]: data.rename(columns= {"Quantity": "Number"})
data
```

Out[78]:

	Price	Quantity	GP%
A	25.0	4.0	0.20
B	50.0	2.0	0.15
C	25.0	5.0	0.20
D	25.0	9.0	0.15
E	15.0	4.0	0.40
F	40.0	6.0	0.15
F	40.0	6.0	0.20
G	25.0	4.0	0.50
H	15.0	5.0	0.20

## 6. Grouping Data

### Groupby

```
In [79]: price= data.groupby("Price")
         type(price)
```

```
Out[79]: pandas.core.groupby.generic.DataFrameGroupBy
```

```
In [80]: price.size() #This is equivalent to data["price"].value_counts
```

```
Out[80]: Price
15.0      2
25.0      4
40.0      2
50.0      1
dtype: int64
```

We can obtain the quantities for various-priced products easily with groupby.

```
In [81]: price.sum()
```

```
Out[81]:
```

	Quantity	GP%
Price		
15.0	9.0	0.60
25.0	22.0	1.05
40.0	12.0	0.35
50.0	2.0	0.15

### Aggregate

Pandas provides the aggregate() method with agg() as alias, which allows for quick summarization and calculation of column data.

```
In [82]: data.groupby("Price")["Quantity"].agg("sum") #This is equivalent to the quan
```

```
Out[82]: Price
15.0      9.0
25.0     22.0
40.0     12.0
50.0      2.0
Name: Quantity, dtype: float64
```

Using aggregate method, we can easily trigger out information.

Such as the median price and quantity among products with different gross profit margins, this could provide some insights for decision-making.

```
In [83]: data.groupby("GP%").agg(["median"])
```

Out[83]:

	Price	Quantity
	median	median
GP%		
0.15	40.0	6.0
0.20	25.0	5.0
0.40	15.0	4.0
0.50	25.0	4.0

## 7. Data Overview

We come back to our df, stock price data.

### Info

Using info() to obtain a brief summary of the DataFrame. It is very convenient for exploratory data analysis.

```
In [84]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 10 entries, 2023-12-08 to 2023-12-17
Data columns (total 3 columns):
#   Column  Non-Null Count  Dtype  
---  -
0    AAPL    10 non-null      float64
1    META    10 non-null      float64
2    NVDA    10 non-null      float64
dtypes: float64(3)
memory usage: 620.0+ bytes
```

### Changing the Datatype

```
In [85]: df["AAPL"] = df["AAPL"].astype(int)
df.head()
```

Out[85]:

	AAPL	META	NVDA
2023-12-08	194	326.59	465.97
2023-12-09	198	330.12	467.13
2023-12-10	199	327.31	471.52
2023-12-11	201	339.51	466.31
2023-12-12	200	336.87	465.19



## Describe

Using describe() to view some basic descriptive statistics details.

```
In [86]: df.describe()
```

Out[86]:


	AAPL	META	NVDA
count	10.000000	10.000000	10.000000
mean	200.300000	336.626000	467.766000
std	2.830391	6.451638	2.117704
min	194.000000	326.590000	465.190000
25%	199.250000	331.807500	466.167500
50%	201.000000	338.620000	467.130000
75%	201.750000	340.275000	468.950000
max	204.000000	345.190000	471.520000

## 8. Basic Visualization With Pandas


Pandas also provides some basic plotting functions to quickly generate simple plots. But for more advanced and complex functionalities, you can use other libraries such as Matplotlib, Seaborn, etc.

```
In [87]: #df.plot(x="Category", y="Value", figsize=(a, b))
df.plot(figsize = (6,4))
# This is equivalent to df.plot.line()
# If not specified arguments, Python will automatically identify and apply t
```

Out[87]: <AxesSubplot:>



```
In [88]: df.plot.bar(figsize=(5,3))
#df.plot.bar()
#df.plot.box()
#df.plot.hist()
#.....
```



## Conclusion

This document introduces many essential but useful techniques in NumPy and Pandas. You now have a basic understanding of data structures such as NumPy Arrays, Series, DataFrame, and have mastered various data manipulation skills.

Equipped with this knowledge, you will be able to quickly get started and demonstrate more efficient workstyle in your data analysis projects. I hope this document helps you solidify your foundation and achieve excellent results by leveraging these 2 libraries.

- Follow my LinkedIn for more information.: <https://www.linkedin.com/in/weihsin-hsu/> (<https://www.linkedin.com/in/weihsin-hsu/>)
- Follow my GitHub for more programming projects.: <https://github.com/endlessnoc> (<https://github.com/endlessnoc>)